

▼ CartPole-v1 with TensorFlow2

Problem gym.openai.com/envs/CartPole-v1/

▼ Setup

```
1 import numpy as np
2 import os
3 import pandas as pd
4 import random
```

```
1 try:
2     import gym
3 except:
4     !pip install gym
5 import gym
6 print(gym.__version__)
```

➞ 0.10.11

```
1 try:
2     import tensorflow as tf
3 except:
4     !pip install tensorflow==2.0.0-alpha0
5 import tensorflow as tf
6 if tf.__version__[0] == "1":
7     !pip install tensorflow==2.0.0-alpha0
8 import tensorflow as tf
9 print(tf.__version__)
```

➞ 2.0.0-alpha0

▼ Visualize gym environment - random action

```
1 import gym
2 env = gym.make("CartPole-v1")
3 no_actions = env.action_space.n
4 no_observations = env.observation_space.shape[0]
5 print(no_actions)
6 print(no_observations)
```

➞ 2
4

Action and observation interpretations:

github.com/openai/gym/blob/master/gym/envs/classic_control/cartpole.py

```
1 observation = env.reset()
```

```

2 for t in range(100):
3     #     env.render()
4     action = env.action_space.sample()
5     observation, reward, done, info = env.step(action)
6     if done:
7         print("Episode finished after {} timesteps".format(t+1))
8         break
9 env.close()

```

➤ Episode finished after 13 timesteps

▼ Train with TF2, Keras

```

1 model = tf.keras.Sequential()
2 model.add(tf.keras.layers.Dense(
3     24,
4     activation="relu",
5     input_shape=(no_observations,)))
6 model.add(tf.keras.layers.Dense(24, activation="relu"))
7 model.add(tf.keras.layers.Dense(no_actions, activation="linear"))
8 model.compile(loss="mse", optimizer=tf.keras.optimizers.Adam(lr=0.0

```

```

1 def get_best_action(state):
2     state_df = pd.DataFrame(columns=["cart_pos", "cart_vel", "pole_
3     state_df.loc[len(state_df),:] = state
4     q = model.predict(state_df.values, batch_size=1)
5     action = pd.Series(q[0]).idxmax()
6     return action, q

```

```

1 GAMMA = 0.95
2 def fit_model(state, action, reward, next_state, done):
3     state_action, state_q = get_best_action(state)
4     next_state_action, next_state_q = get_best_action(next_state)
5     q_update = reward
6     if not done:
7         q_update = (reward + GAMMA * next_state_q[0][next_state_act
8     state_q[0][action] = q_update
9     state_df = pd.DataFrame(columns=["cart_pos", "cart_vel", "pole_
10    state_df.loc[len(state_df),:] = state
11    model.fit(state_df.values, state_q, epochs=1, batch_size=1)

```

```

1 %%capture
2 cap_T = 1500
3 for ep in range(cap_T):
4     state = env.reset()
5     for t in range(500):
6         #     env.render()
7         if random.random() < (1 - ep / cap_T):
8             action = env.action_space.sample()
9         else:
10            action, q = get_best_action(state)
11            next_state, reward, done, info = env.step(action)
12            fit_model(state, action, reward, next_state, done)
13            state = next_state
14            if done:
15                print("Episode finished after {} timesteps".format(t+1))
16                break

```

▼ Test with TF2, Keras

```
1 # env.render()
2 for ep in range(10):
3     state = env.reset()
4     for t in range(1000):
5         action, q = get_best_action(state)
6         next_state, reward, done, info = env.step(action)
7         if done:
8             print("Episode {} finished after {} timesteps".format(e
9                 break
10         state = next_state
11 env.close()
```

```
➤ Episode 0 finished after 99 timesteps
Episode 1 finished after 101 timesteps
Episode 2 finished after 103 timesteps
Episode 3 finished after 104 timesteps
Episode 4 finished after 99 timesteps
Episode 5 finished after 103 timesteps
Episode 6 finished after 104 timesteps
Episode 7 finished after 106 timesteps
Episode 8 finished after 102 timesteps
Episode 9 finished after 105 timesteps
```